
Open file format for MR sequences

Version 1.1.0

Kelvin Layton
Maxim Zaitsev

University Medical Centre Freiburg
`kelvin.layton@uniklinik-freiburg.de`
`maxim.zaitsev@uniklinik-freiburg.de`

This file specification is part of the Pulseseq project:



This file format is licensed under the Creative Commons Attribution 4.0
International License. To view a copy of this license, visit
<http://creativecommons.org/licenses/by/4.0/>

Contents

1	Introduction	2
1.1	Example	3
2	Specification	4
2.1	Overall Description	4
2.2	Identification numbers	4
2.3	General	5
2.4	Version	6
2.5	Blocks	6
2.6	Events	7
2.6.1	RF	7
2.6.2	Gradients	8
2.6.3	ADC	9
2.6.4	Delays	9
2.7	Shapes	10
2.7.1	Compression	10
3	Binary files	11
3.1	File and section codes	12
3.2	Section details	12
4	Source code	13
5	Examples	14
5.1	Free induction decay	14
5.2	Point-resolved spectroscopy (PRESS)	15
5.3	Gradient echo	16

Revision History

Version	Date	Description
1.0	26 Jun 2014	Draft specification
1.01	11 Jun 2015	Included draft of binary specification
1.1.0	11 Jul 2017	Changed versioning scheme

1 Introduction

The purpose of this file format is to compactly represent a magnetic resonance (MR) sequence in an open and vendor independent manner. Sequences for both NMR spectroscopy and MRI can be represented. The format is intended for research and educational purposes and currently omits complex sequence features such as physiological triggering or logical coordinate-frame transformations. The format has been developed with the following **design goals**:

1. **Human-readable:** The basic sequence structure should be easily understood without processing.
2. **Easily parsed:** The format should be easy for a computer program to parse without the need for external libraries.
3. **Vendor independent:** The sequence format must not contain definitions specific to a particular hardware manufacturer.
4. **Compact:** The sequence must not contain redundant definitions. As such, re-use of existing definitions at different times is inherent in the sequence format.
5. **Low-level:** The format should be sufficiently low-level. This allows for maximum flexibility of sequence specifications and does not limit the high-level design tools.

The first goal of human readability necessitates a text-file format. The file format is intended to describe a sequence that can be run on scanner hardware. Therefore, the second goal of machine readability ensures that the file can be read and interpreted without the use of external libraries that might be unavailable on different platforms. This prohibits the use of an existing markup language like XML, which are not straightforward to parse. Further, units that are inherently hardware dependent have been avoided, such as ‘Volts’ or ‘Tesla’.

1.1 Example

Before defining the detailed specification, a simple example is presented to demonstrate the main structure of a sequence. Below is a simple FID experiment with RF pulse, delay and ADC readout.

```
# Sequence Blocks
# Created by JEMRIS 2.8

[VERSION]
major 1
minor 1
revision 0

[DEFINITIONS]
Scan.ID 2
Num.Blocks 3

# Format of blocks:
## D RF GX GY GZ ADC
[BLOCKS]
1 0 1 0 0 0 0
2 1 0 0 0 0 0
3 0 0 0 0 0 1

# Format of RF events:
# id amplitude mag.id phase.id freq phase
# .. Hz .... Hz rad
[RF]
1 2500 1 2 0 0

# Format of ADC events:
# id num dwell delay freq phase
# .. .. ns us Hz rad
[ADC]
1 64 50000 0 0 0

# Format of delays:
# id delay (us)
[DELAYS]
1 5000

# Sequence Shapes

[SHAPES]

shape_id 1
num.samples 100
1
0
0
97

shape_id 2
num.samples 100
```

```

0
0
98

```

2 Specification

2.1 Overall Description




Figure 1: Visualisation of the hierarchical structure of the file format for an FID example.

The sequence description consists of a three-level hierarchical structure as demonstrated in Figure 1. This ensures the most compact representation, since repeating events (or shapes) are only declared once.

Blocks At the top level the sequence is specified by a number of blocks. Each block refers to one or more event(s) to be executed simultaneously.

Events The definition of events are dependent on the type (e.g. gradient, ADC) but some types may refer to one or more basic shapes.

Shapes A shape is a compressed list of samples. The uncompressed samples can describe, for example, the RF pulse shape or an arbitrary gradient shape. The compression scheme is a type of run-length encoding.

2.2 Identification numbers

A key idea of the hierarchical sequence structure is the use of ID numbers to refer to objects defined one level below in the hierarchy. For example, blocks contain the ID of one or more event(s). Likewise, events may contain the ID of one or more shape(s). Restrictions are placed on these IDs to ensure consistency:

- IDs must be positive integers.
- The ID of each shape must be unique.
- The ID of events within a single class must be unique. For example, an RF event and delay event may both use an ID of 1, since the events are of a different class. An exception is the [GRADIENTS] and [TRAP] events where the union of these sets must contain unique event IDs.
- There are no restrictions on the ordering of IDs, although sequential ordering is often implemented.
- There are no restrictions on the first ID of an event class.

2.3 General

Comments are specified by a line starting with a hash.

```
# This is a comment
```

The sequence may contain general definitions indicated by the [DEFINITIONS] keyword. Each definition is defined on a new line with the following format

```
<key> <value>
```

This defines a list of user-specified key/value pairs.

Tag	Type	Description	Units
<key>	text	Name of user definition	–
<value>	any	Value of definition	–

Example:

```
[DEFINITIONS]
Scan.ID 237
FOV 150 150
TE 5
TR 100
Num.Blocks 3072
```

It is important to note that all definitions are optional and do not affect the basic sequence execution. Precise timing is given by the low-level specification of events. The definitions section may be used for user-specific purposes, including attaching metadata or hardware-dependent parameters.

2.4 Version

The versioning scheme is $\langle\text{major}\rangle.\langle\text{minor}\rangle.\langle\text{revision}\rangle$.

```
[VERSION]
major <major>
minor <minor>
revision <revision>
```

Providing the version of the standard is vital for sequence execution. If this section is not provided, the interpreter sequence may either refuse execution or assume version 1.0.0.

Example:

```
[VERSION]
major 1
minor 1
revision 0
```

2.5 Blocks

The section containing sequence blocks is declared with the [BLOCKS] keyword. Each subsequence line declares a single block specified by a list of six event IDs. An ID of 0 indicated no event.

```
<id> <delay> <rf> <gx> <gy> <gz> <adc>
```

Each tag represents the ID of the corresponding event.

Tag	Type	Description	Units
<id>	integer	ID of the sequence block	—
<delay>	integer	ID of the delay event	—
<rf>	integer	ID of the RF event	—
<gx>	integer	ID of the gradient event on the X channel	—
<gy>	integer	ID of the gradient event on the Y channel	—
<gz>	integer	ID of the gradient event on the Z channel	—
<adc>	integer	ID of the ADC event	—

The sequence must declare at least one block. Any non-zero number of blocks may be defined. The blocks are executed sequentially. The duration of each block is defined by the longest event comprising this block. X, Y and Z refer to physical scanner gradient channels.

Example:

```
[BLOCKS]
  1  0  1  0  0  2  0
```

The block above is the first in the sequence and contains an RF pulse with ID of 1 and a z-gradient pulse with ID of 2. The block has no delay, X gradient, Y gradient or ADC events, indicated by zero IDs.

2.6 Events

Events are defined in sections, each starting with one the following keywords: [RF], [GRADIENTS], [TRAP], [ADC] or [DELAYS]. Each event is specified on a single line and contains an ID followed by type-specific definition.

2.6.1 RF

The RF section is declared with the [RF] keyword. Following this declaration, each RF event is specified by a single line containing six numbers.

```
<id> <amp> <mag_id> <phase_id> <freq> <phase>
```

The specifiers are

Tag	Type	Description	Units
<id>	integer	ID of the RF event	–
<amp>	float	Peak amplitude	Hz
<mag_id>	integer	Shape ID for magnitude profile	–
<phase_id>	integer	Shape ID for phase profile	–
<freq>	float	Frequency offset	Hz
<phase>	float	Phase offset	rad

Example:

```
[RF]
  1  2500  1  2  0.000  0.000
```


In the example above, the RF pulse has ID of 1, peak amplitude of 2500Hz. The magnitude and phase profiles are defined with the shapes of ID 1 and 2, respectively. The RF pulse does not have any frequency or phase offset.

2.6.2 Gradients

Gradient events are declared in two sections. Arbitrary gradients are in a section declared with the **[GRADIENTS]** keyword. Each line in the section is an arbitrary gradient specified by three numbers,

```
<id> <amp> <shape_id>
```

Trapezoidal gradients are in a section declared with the **[TRAP]** keyword. Each line in the section is a trapezoidal gradients specified by five numbers,

```
<id> <amp> <rise> <flat> <fall>
```

The specifiers are

Tag	Type	Description	Units
<id>	integer	ID of the gradient event	–
<amp>	float	Peak amplitude	Hz/m
<shape_id>	integer	Shape ID for arbitrary gradient waveform	–
<rise>	integer	Rise time of the trapezoid	μs
<flat>	integer	Flat-top time of the trapezoid	μs
<fall>	integer	Fall time of the trapezoid	μs

The gradient ID must be unique across both arbitrary and trapezoid gradients. That is, a trapezoid gradient cannot have the same ID as an arbitrary gradient.

Example:

```
[GRADIENTS]
 1 159154.9 3
[TRAP]
 2 25000 30 940 30
 3 20066.89 10 980 10
```

The example above contains two gradients: an arbitrary gradient with peak amplitude of approximately 159kHz/m and shape ID 3 and two trapezoid gradients (IDs 2 and 3) with duration 1ms specified by amplitude, rise time, flat-top time and fall time.

2.6.3 ADC

The ADC section is declared with the [ADC] keyword. Each line in the section is an ADC event specified by six numbers,

```
<id> <num> <dwel1> <delay> <freq> <phase>
```

The specifiers are

Tag	Type	Description	Units
<id>	integer	ID of the ADC event	–
<num>	integer	Number of samples	–
<dwel1>	float	The ADC dwell time	ns
<delay>	integer	Delay between start of block and first sample	μ s
<freq>	float	Frequency offset of ADC receiver	Hz
<phase>	float	Phase offset of ADC receiver	rad

The duration of the ADC readout is given by the product of <num> and <dwel1>.

Example:

```
[ADC]
1      512  5000  0  0.000  0.000
```

The example above contains an ADC with 512 samples, and dwell time of 5000ns, and no frequency and phase offsets. The frequency and phase offset are used, for example, for RF spoiling or inplane shifting of the FOV.

2.6.4 Delays

The delay section is declared with the [DELAYS] keyword. Each line in the section is a delay event specified by two numbers,

```
<id> <delay>
```

The specifiers are

Tag	Type	Description	Units
<id>	integer	ID of the delay event	–
<delay>	integer	Delay before starting the execution of the block	μ s

2.7 Shapes

The shape section is declared with the [SHAPES] keyword. Each shape is then declared with a header followed by a list of samples values (one per line). The end of the shape definition is declared with a blank line.

```
Shape.ID <id>
Num.Uncompressed <num>
<sample_1>
<sample_2>
...
```

The specifiers are

Tag	Type	Description	Units
<id>	integer	ID of the shape	–
<num>	integer	Number of samples of the uncompressed shape	–
<sample_n>	integer	The n^{th} sample of the compressed shape	–

The decompressed samples must be in the normalised range $[0, 1]$. Since the purpose of this section is to define the basic shape of a gradient or RF pulse, the amplitude information is defined in the events section. This allows the same shape to be used with different amplitudes, without repeated definitions.

The number of points after decompressing all samples defined in a shape must equal the number declared in <num>.

2.7.1 Compression

Storing every sample of the shape would lead to very large sequence descriptions. Suppose a sequence contains a block RF pulse for 4ms and a sinusoidally-ramped constant gradient for 100ms. Assuming sampling times of $1\mu\text{s}$ and $10\mu\text{s}$ for the RF and gradients, respectively, 14000 samples would be required. Instead, the shapes are compressed by **encoding the derivative in a run-length compressed** format.

Example 1: A shape consisting of a ramp-up, constant and ramp-down is encoded as follows

Shape		Step 1 (derivative)		Step 2 (compression)
0.0		0.0		0
0.1		0.1		0.1
0.25		0.15		0.15
0.5		0.25		0.25
1.0		0.5		0.5
1.0		0.0		0.0
1.0	→	0.0	→	0.0
1.0		0.0		4
1.0		0.0		-0.25
1.0		0.0		-0.25
1.0		0.0		2
0.75		-0.25		
0.5		-0.25		
0.25		-0.25		
0.0		-0.25		

Example 2: A shape with 100 zeros values

Shape		Step 1 (derivative)		Step 2 (compression)
0.0		0.0		0
0.0	→	0.0	→	0
...		...		98
0.0		0.0		

Example 3: A shape with a constant value of 1.0 for 100 samples

Shape		Step 1 (derivative)		Step 2 (compression)
1.0		1.0		1.0
1.0	→	0.0	→	0
...		...		0
1.0		0.0		97

3 Binary files

The specification described in Section 2 can be implemented as a binary file. The same general principles apply with specific modifications outlined here. The basic structure of a binary *pulseq* file is depicted below,

0x01	p	u	l	s	e	q	0x02
version major							
version minor							
version revision							
section code							
number of events							
data							
section code							
number of events							
data							
...							

3.1 File and section codes

A binary *Pulseq* file begins with the 64 bit code 0x0170756C73657102 (the characters **pulseq** enclosed by 0x01 and 0x02) followed by three integers describing the file version (major, minor, revision). The remaining file is made up of multiple sections each with an integer section code followed by section-specific storage. The section codes corresponding to text file tags are

Section	Section code
[DEFINITIONS]	0xFFFFFFFF 0x00000001
[BLOCKS]	0xFFFFFFFF 0x00000002
[RF]	0xFFFFFFFF 0x00000003
[GRADIENTS]	0xFFFFFFFF 0x00000004
[TRAP]	0xFFFFFFFF 0x00000005
[ADC]	0xFFFFFFFF 0x00000006
[DELAYS]	0xFFFFFFFF 0x00000007
[SHAPES]	0xFFFFFFFF 0x00000008

3.2 Section details

The number of entries in each section is written as an integer directly after the section code.

Unlike the text format, the [DEFINITIONS] section is declared with the section code followed by *four* fields per definition

```
<key> <type> <num> <value_1> <value_2>...
```

The **<key>** is a null-terminated ASCII encoded string. **<type>** is an integer code either 1 to denote integer or 2 to denote floating-point. **<num>** is an integer defining the number of values in this definition and **<value_n>** are the values of given type. The next definition immediately follows otherwise a new section code follows or it is the end of the file.

The [SHAPES] section is declared with the appropriate section code followed by

```
<id> <num_uncompressed> <num_compressed> <sample_1> <sample_2>...
```

The **<id>** is an integer shape ID, **<num_uncompressed>** is the number of samples after decompression, **<num_compressed>** is the number of compressed samples for this shape. Following are exactly the specified number of compressed samples stored as floating-point numbers. The next shape begins immediately after the last otherwise a new section code is declared (or end-of-file) to indicate the end of the [SHAPES] section.

The remaining sections are declared exactly as in Section 2 with storage defined by the value type. In general, **integer** types are stored as 64-bit unsigned integers while **floating-point** numbers are in the 64-bit IEEE 754 floating-point format. Exceptions are **<id>** values, which are 32-bit integers; and **<type>** and **<num>** values of the [DEFINITIONS] entries, which are 8-bit integers.

4 Source code

This specification is distributed with source code for reading and writing the sequences file format described here. MATLAB code is provided for detailed sequence generation, visualisation, as well as reading and writing sequence files. A C++ class and example program is also provided for reading sequence files. Detailed documentation and latest updates of this code are available here: <http://pulseseq.github.io/>.

5 Examples

5.1 Free induction decay

```
# Sequence Blocks
# Created by JEMRIS 2.8

[VERSION]
major 1
minor 1
revision 0

[DEFINITIONS]
Scan.ID 2
Num.Blocks 3

# Format of blocks:
## D RF GX GY GZ ADC
[BLOCKS]
1 0 1 0 0 0 0
2 1 0 0 0 0 0
3 0 0 0 0 0 1

# Format of RF events:
# id amplitude mag.id phase.id freq phase
# .. Hz .... Hz rad
[RF]
1 2500 1 2 0 0

# Format of ADC events:
# id num dwell delay freq phase
# .. ns us Hz rad
[ADC]
1 64 50000 0 0 0

# Format of delays:
# id delay (us)
[DELAYS]
1 5000

# Sequence Shapes

[SHAPES]

shape_id 1
num.samples 100
1
0
0
97

shape_id 2
num.samples 100
0
0
```

5.2 Point-resolved spectroscopy (PRESS)

```

# Sequence Blocks
# Created by JEMRIS 2.8

[DEFINITIONS]
Scan_ID 3
Num.Blocks 8

# Format of blocks:
## D RF GX GY GZ ADC
[BLOCKS]
1 0 1 0 1 0 0
2 0 0 0 2 0 0
3 1 0 0 0 0 0
4 0 2 3 0 0 0
5 2 0 0 0 0 0
6 0 2 0 0 4 0
7 3 0 0 0 0 0
8 0 0 0 0 0 1

# Format of RF events:
# id amplitude mag_id phase_id freq phase
# .. Hz .... Hz rad
[RF]
1 246.863892 1 2 0 0
2 493.727784 1 2 0 0

# Format of trapezoid gradients:
# id amplitude rise flat fall
# .. Hz/m us us us
[TRAP]
1 200106.3 30 3940 30
2 -201629.9 30 1940 30
3 200106.3 30 3940 30
4 200106.3 30 3940 30

# Format of ADC events:
# id num dwell delay freq phase
# .. .. ns us Hz rad
[ADC]
1 256 25000 0 0 0

# Format of delays:
# id delay (us)
[DELAYS]
1 19000
2 46000
3 44800

# Sequence Shapes

```



```

[SHAPES]

shape_id 1
num.samples 4000
3.856279e-11
3.588134e-08
3.588134e-08
8
2.323701e-07
2.323701e-07
8
6.188511e-07
6.188511e-07
8
1.197773e-06
1.197773e-06
8
1.970951e-06
1.970951e-06
8
2.939553e-06
2.939553e-06
8
4.104086e-06
4.104086e-06
8
5.464388e-06
5.464388e-06
8
<<truncated>>

shape_id 2
num.samples 4000
0.5
0
0
997
-0.5
0
0
1997
0.5
0
0
997

```

5.3 Gradient echo

```

# Sequence Blocks
# Created by JEMRIS 2.8

[DEFINITIONS]
Scan.ID 2
Num.Blocks 160

```

```

# Format of blocks:
# # D RF GX GY GZ ADC
[BLOCKS]
1 0 1 0 0 1 0
2 0 0 2 3 4 0
3 1 0 0 0 0 0
4 0 0 5 0 0 1
5 2 0 0 0 0 0
6 0 1 0 0 1 0
7 0 0 2 6 4 0
8 1 0 0 0 0 0
9 0 0 5 0 0 1
10 2 0 0 0 0 0
11 0 1 0 0 1 0
12 0 0 2 7 4 0
<<truncated>>

# Format of RF events:
# id amplitude mag_id phase_id freq phase
# .. Hz .... Hz rad
[RF]
1 276.908986 1 2 0 0

# Format of trapezoid gradients:
# id amplitude rise flat fall
# .. Hz/m us us us
[TRAP]
1 159154.9 30 4000 30
2 -26797.66 10 2980 10
3 -26755.85 10 2980 10
4 -107616.5 20 2960 20
5 25000 10 6400 10
6 -25083.61 10 2980 10
7 -23411.37 10 2980 10
<<truncated>>

# Format of ADC events:
# id num dwell delay freq phase
# .. .. ns us Hz rad
[ADC]
1 32 200000 10 0 0

# Format of delays:
# id delay (us)
[DELAYS]
1 1760
2 184770

# Sequence Shapes
[SHAPES]

shape_id 1
num_samples 4000
0.0002500625
0.0005026691

```

```
0.0005026691
8
0.0005072244
0.0005072244
8
0.0005113487
0.0005113487
8
0.0005150341
0.0005150341
8
0.0005182732
0.0005182732
<<truncated>>

shape_id 2
num_samples 4000
0.5
0
0
997
-0.5
0
0
1997
0.5
0
0
997
```